

```
In [1]: import pandas as pd
import numpy as np
from matplotlib import pyplot as plt
import seaborn as sns
from sklearn.linear_model import LinearRegression, LogisticRegression, Lasso, Ridge, ElasticNet
from sklearn.model_selection import train_test_split
```

```
In [2]: df=pd.read_csv("C:/Users/user/Downloads/FP1_air/csvs_per_year/csvs_per_year/madrid_2003")
df
```

Out[2]:

	date	BEN	CO	EBE	MXV	NMHC	NO_2	NOx	OXY	O_3	PM10	PXY
0	2003-03-01 01:00:00	NaN	1.72	NaN	NaN	NaN	73.900002	316.299988	NaN	10.550000	55.209999	NaN
1	2003-03-01 01:00:00	NaN	1.45	NaN	NaN	0.26	72.110001	250.000000	0.73	6.720000	52.389999	NaN
2	2003-03-01 01:00:00	NaN	1.57	NaN	NaN	NaN	80.559998	224.199997	NaN	21.049999	63.240002	NaN
3	2003-03-01 01:00:00	NaN	2.45	NaN	NaN	NaN	78.370003	450.399994	NaN	4.220000	67.839996	NaN
4	2003-03-01 01:00:00	NaN	3.26	NaN	NaN	NaN	96.250000	479.100006	NaN	8.460000	95.779999	NaN
...	...	...	...	...	...	...	...	...	...	...	...	...
243979	2003-10-01 00:00:00	0.20	0.16	2.01	3.17	0.02	31.799999	32.299999	1.68	34.049999	7.380000	1.20
243980	2003-10-01 00:00:00	0.32	0.08	0.36	0.72	NaN	10.450000	14.760000	1.00	34.610001	7.400000	0.50
243981	2003-10-01 00:00:00	NaN	NaN	NaN	NaN	0.07	34.639999	50.810001	NaN	32.160000	16.830000	NaN
243982	2003-10-01 00:00:00	NaN	NaN	NaN	NaN	0.07	32.580002	41.020000	NaN	NaN	13.570000	NaN
243983	2003-10-01 00:00:00	1.00	0.29	2.15	6.41	0.07	37.150002	56.849998	2.28	21.480000	12.350000	2.43

243984 rows × 16 columns

```
In [3]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 243984 entries, 0 to 243983
Data columns (total 16 columns):
 #   Column      Non-Null Count  Dtype  
---  -
 0   date        243984 non-null object  
 1   BEN         69745 non-null float64
 2   CO          225340 non-null float64
 3   EBE         61244 non-null float64
 4   MXY         42045 non-null float64
 5   NMHC        111951 non-null float64
 6   NO_2        242625 non-null float64
 7   NOx         242629 non-null float64
 8   OXY         42072 non-null float64
 9   O_3         234131 non-null float64
10  PM10        240896 non-null float64
11  PXY         42063 non-null float64
12  SO_2        242729 non-null float64
13  TCH         111991 non-null float64
14  TOL         69439 non-null float64
15  station     243984 non-null int64  
dtypes: float64(14), int64(1), object(1)
memory usage: 29.8+ MB
```

```
In [4]: df1=df.dropna()
df1
```

Out[4]:

	date	BEN	CO	EBE	MXV	NMHC	NO_2	NOx	OXY	O_3	PM10	PXY
5	2003-03-01 01:00:00	8.41	1.94	9.83	21.49	0.45	90.300003	384.899994	9.48	9.950000	95.150002	7.94
23	2003-03-01 01:00:00	3.46	1.27	3.43	7.08	0.18	54.250000	173.300003	3.37	6.540000	53.009998	2.62
27	2003-03-01 01:00:00	6.39	1.79	5.75	10.88	0.33	75.459999	281.100006	3.68	6.690000	63.840000	4.24
33	2003-03-01 02:00:00	7.42	1.47	10.63	24.73	0.35	83.309998	277.200012	11.00	9.900000	58.880001	8.93
51	2003-03-01 02:00:00	3.62	1.29	3.20	7.08	0.19	42.209999	166.300003	3.41	6.380000	47.599998	2.70
...	...	...	...	...	...	...	...	...	...	...	...	...
243955	2003-09-30 23:00:00	1.75	0.41	3.07	9.38	0.09	46.290001	77.709999	3.11	18.280001	7.520000	3.48
243957	2003-10-01 00:00:00	2.35	0.60	3.88	10.86	0.11	61.240002	133.100006	0.89	10.900000	10.240000	3.89
243961	2003-10-01 00:00:00	2.97	0.82	4.53	10.88	0.05	36.529999	131.300003	5.52	12.940000	25.680000	4.13
243979	2003-10-01 00:00:00	0.20	0.16	2.01	3.17	0.02	31.799999	32.299999	1.68	34.049999	7.380000	1.20
243983	2003-10-01 00:00:00	1.00	0.29	2.15	6.41	0.07	37.150002	56.849998	2.28	21.480000	12.350000	2.43


33010 rows × 16 columns



```
In [5]: df1=df1.drop(["date"],axis=1)
```


```
In [6]: sns.heatmap(df1.corr())
```

```
Out[6]: <AxesSubplot:>
```



```
In [7]: plt.plot(df1["EBE"],df1["PXY"],"o")
```

```
Out[7]: [<matplotlib.lines.Line2D at 0x255505b0490>]
```



```
In [8]: data=df[["EBE","PXY"]]
```

```
In [9]: # sns.stripplot(x=df["EBE"],y=df["PXY"],jitter=True,marker='o',color='blue')
```

```
In [10]: x=df1.drop(["EBE"],axis=1)
y=df1["EBE"]
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```


## Linear

```
In [11]: li=LinearRegression()
li.fit(x_train,y_train)
```

```
Out[11]: LinearRegression()
```

```
In [12]: prediction=li.predict(x_test)
plt.scatter(y_test,prediction)
```

```
Out[12]: <matplotlib.collections.PathCollection at 0x25550a8dee0>
```



```
In [13]: lis=li.score(x_test,y_test)
```

```
In [14]: df1["TCH"].value_counts()
```

```
Out[14]: 1.30    1344
         1.31    1342
         1.32    1281
         1.27    1279
         1.29    1262
         ...
         3.50      1
         3.87      1
         3.21      1
         3.14      1
         1.01      1
         Name: TCH, Length: 243, dtype: int64
```

```
In [15]: df1.loc[df1["TCH"]<1.40,"TCH"]=1
df1.loc[df1["TCH"]>1.40,"TCH"]=2
df1["TCH"].value_counts()
```

```
Out[15]: 1.0    21614
         2.0    11396
         Name: TCH, dtype: int64
```


## Lasso

```
In [16]: la=Lasso(alpha=5)
la.fit(x_train,y_train)
```

```
Out[16]: Lasso(alpha=5)
```

```
In [17]: prediction1=la.predict(x_test)
plt.scatter(y_test,prediction1)
```

```
Out[17]: <matplotlib.collections.PathCollection at 0x25550acbc70>
```



```
In [18]: las=la.score(x_test,y_test)
```


## Ridge

```
In [19]: rr=Ridge(alpha=1)
rr.fit(x_train,y_train)
```

```
Out[19]: Ridge(alpha=1)
```

```
In [20]: prediction2=rr.predict(x_test)
plt.scatter(y_test,prediction2)
```

```
Out[20]: <matplotlib.collections.PathCollection at 0x255504ca8e0>
```



```
In [21]: rrs=rr.score(x_test,y_test)
```


## ElasticNet

```
In [22]: en=ElasticNet()
en.fit(x_train,y_train)
```

```
Out[22]: ElasticNet()
```

```
In [23]: prediction2=rr.predict(x_test)
plt.scatter(y_test,prediction2)
```

```
Out[23]: <matplotlib.collections.PathCollection at 0x25551ce3cd0>
```



```
In [24]: ens=en.score(x_test,y_test)
```

```
In [25]: print(rr.score(x_test,y_test))
rr.score(x_train,y_train)
```

```
0.9195176980660054
```

```
Out[25]: 0.9131593670000937
```

## Logistic

```
In [26]: g={"TCH":{1.0:"Low",2.0:"High"}}
df1=df1.replace(g)
df1["TCH"].value_counts()
```

```
Out[26]: Low      21614
High      11396
Name: TCH, dtype: int64
```

```
In [27]: x=df1.drop(["TCH"],axis=1)
y=df1["TCH"]
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

```
In [28]: lo=LogisticRegression()
lo.fit(x_train,y_train)
```

```
Out[28]: LogisticRegression()
```

```
In [29]: prediction3=lo.predict(x_test)
plt.scatter(y_test,prediction3)
```

```
Out[29]: <matplotlib.collections.PathCollection at 0x25551986520>
```



```
In [30]: los=lo.score(x_test,y_test)
```

## Random Forest

```
In [31]: from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV
```

```
In [32]: g1={"TCH":{"Low":1.0,"High":2.0}}
df1=df1.replace(g1)
```

```
In [33]: x=df1.drop(["TCH"],axis=1)
y=df1["TCH"]
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

```
In [34]: rfc=RandomForestClassifier()
rfc.fit(x_train,y_train)
```

```
Out[34]: RandomForestClassifier()
```

```
In [35]: parameter={
    'max_depth':[1,2,4,5,6],
    'min_samples_leaf':[5,10,15,20,25],
    'n_estimators':[10,20,30,40,50]
}
```

```
In [36]: grid_search=GridSearchCV(estimator=rfc,param_grid=parameter,cv=2,scoring="accuracy")
grid_search.fit(x_train,y_train)
```

```
Out[36]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
    param_grid={'max_depth': [1, 2, 4, 5, 6],
    'min_samples_leaf': [5, 10, 15, 20, 25],
    'n_estimators': [10, 20, 30, 40, 50]},
    scoring='accuracy')
```



```
In [37]: rfcs=grid_search.best_score_
```

```
In [38]: rfc_best=grid_search.best_estimator_
```

```
In [39]: from sklearn.tree import plot_tree

plt.figure(figsize=(80,40))
plot_tree(rfc_best.estimators_[5],feature_names=x.columns,class_names=['Yes','No'],fill
Text(1075.9999999999998, 1597.8285714285716, 'NMHC <= 0.155\ngini = 0.497\nsamples
= 1496\nvalue = [1278, 1106]\nclass = Yes'),
Text(1355.142857142857, 1087.2, 'OXY <= 3.16\ngini = 0.296\nsamples = 845\nvalue =
[1080, 238]\nclass = Yes'),
Text(1195.7142857142856, 776.5714285714287, 'NOx <= 101.4\ngini = 0.418\nsamples =
314\nvalue = [350, 148]\nclass = Yes'),
Text(1116.0, 465.9428571428573, 'O_3 <= 13.42\ngini = 0.309\nsamples = 186\nvalue =
[241, 57]\nclass = Yes'),
Text(1076.142857142857, 155.3142857142857, 'gini = 0.481\nsamples = 53\nvalue = [5
5, 37]\nclass = Yes'),
Text(1155.8571428571427, 155.3142857142857, 'gini = 0.175\nsamples = 133\nvalue =
[186, 20]\nclass = Yes'),
Text(1275.4285714285713, 465.9428571428573, 'station <= 28079015.0\ngini = 0.496\ns
amples = 128\nvalue = [109, 91]\nclass = Yes'),
Text(1235.5714285714284, 155.3142857142857, 'gini = 0.26\nsamples = 31\nvalue = [4
4, 8]\nclass = Yes'),
Text(1315.2857142857142, 155.3142857142857, 'gini = 0.493\nsamples = 97\nvalue = [6
5, 83]\nclass = No'),
Text(1514.5714285714284, 776.5714285714287, 'O_3 <= 13.285\ngini = 0.195\nsamples =
531\nvalue = [730, 90]\nclass = Yes'),
```

```
In [40]: print("Linear:",lis)
print("Lasso:",las)
print("Ridge:",rrs)
print("ElasticNet:",ens)
print("Logistic:",los)
print("Random Forest:",rfcs)
```

```
Linear: 0.9195103306695078
Lasso: 0.6724055165668984
Ridge: 0.9195176980660054
ElasticNet: 0.8935055393696031
Logistic: 0.6567706755528627
Random Forest: 0.8821569724921972
```

## Best Model is Ridge Regression

```
In [41]: df2=pd.read_csv("C:/Users/user/Downloads/FP1_air/csvs_per_year/csvs_per_year/madrid_2004-04-01.csv")
df2
```

Out[41]:

	date	BEN	CO	EBE	MXV	NMHC	NO_2	NOx	OXY	O_3	PM10	P
0	2004-08-01 01:00:00	NaN	0.66	NaN	NaN	NaN	89.550003	118.900002	NaN	40.020000	39.990002	25.860000
1	2004-08-01 01:00:00	2.66	0.54	2.99	6.08	0.18	51.799999	53.860001	3.28	51.689999	22.950001	12.040000
2	2004-08-01 01:00:00	NaN	1.02	NaN	NaN	NaN	93.389999	138.600006	NaN	20.860001	49.480000	12.040000
3	2004-08-01 01:00:00	NaN	0.53	NaN	NaN	NaN	87.290001	105.000000	NaN	36.730000	31.070000	12.040000
4	2004-08-01 01:00:00	NaN	0.17	NaN	NaN	NaN	34.910000	35.349998	NaN	86.269997	54.080002	12.040000
...	...	...	...	...	...	...	...	...	...	...	...	...
245491	2004-06-01 00:00:00	0.75	0.21	0.85	1.55	0.07	59.580002	64.389999	0.66	33.029999	30.900000	14.860000
245492	2004-06-01 00:00:00	2.49	0.75	2.44	4.57	NaN	97.139999	146.899994	2.34	7.740000	37.689999	12.040000
245493	2004-06-01 00:00:00	NaN	NaN	NaN	NaN	0.13	102.699997	132.600006	NaN	17.809999	22.840000	12.040000
245494	2004-06-01 00:00:00	NaN	NaN	NaN	NaN	0.09	82.599998	102.599998	NaN	NaN	45.630001	12.040000
245495	2004-06-01 00:00:00	3.01	0.67	2.78	5.12	0.20	92.550003	141.000000	2.60	11.460000	24.389999	17.950000

245496 rows × 17 columns



```
In [42]: df2.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 245496 entries, 0 to 245495
Data columns (total 17 columns):
 #   Column      Non-Null Count  Dtype  
---  -
 0   date        245496 non-null object  
 1   BEN         65158 non-null float64
 2   CO          226043 non-null float64
 3   EBE         56781 non-null float64
 4   MXY         39867 non-null float64
 5   NMHC        107630 non-null float64
 6   NO_2        243280 non-null float64
 7   NOx         243283 non-null float64
 8   OXY         39882 non-null float64
 9   O_3         233811 non-null float64
10  PM10        234655 non-null float64
11  PM25        58145 non-null float64
12  PXY         39891 non-null float64
13  SO_2        243402 non-null float64
14  TCH         107650 non-null float64
15  TOL         64914 non-null float64
16  station     245496 non-null int64  
dtypes: float64(15), int64(1), object(1)
memory usage: 31.8+ MB
```

```
In [43]: df3=df2.dropna()  
df3
```

Out[43]:

	date	BEN	CO	EBE	MXV	NMHC	NO_2	NOx	OXY	O_3	PM10	P
5	2004-08-01 01:00:00	3.24	0.63	5.55	9.72	0.06	103.800003	144.800003	5.04	32.480000	59.110001	38.049
22	2004-08-01 01:00:00	0.55	0.36	0.54	0.86	0.07	31.980000	32.799999	0.50	79.040001	43.549999	22.780
26	2004-08-01 01:00:00	1.80	0.46	2.28	4.62	0.21	62.259998	75.470001	2.47	54.419998	46.630001	29.459
32	2004-08-01 02:00:00	1.94	0.67	3.14	4.91	0.06	113.500000	165.800003	2.56	26.980000	86.930000	45.639
49	2004-08-01 02:00:00	0.29	0.30	0.47	0.76	0.07	33.919998	34.840000	0.46	75.570000	48.959999	25.959
...	...	...	...	...	...	...	...	...	...	...	...	...
245463	2004-05-31 23:00:00	0.62	0.08	0.54	0.70	0.04	44.360001	45.450001	0.42	43.419998	19.290001	13.190
245467	2004-05-31 23:00:00	2.39	0.67	2.49	3.92	0.20	89.809998	132.800003	2.09	14.740000	31.809999	20.370
245473	2004-06-01 00:00:00	3.72	1.12	4.33	8.79	0.24	113.900002	253.600006	4.51	9.380000	21.219999	16.660
245491	2004-06-01 00:00:00	0.75	0.21	0.85	1.55	0.07	59.580002	64.389999	0.66	33.029999	30.900000	14.860
245495	2004-06-01 00:00:00	3.01	0.67	2.78	5.12	0.20	92.550003	141.000000	2.60	11.460000	24.389999	17.959


19397 rows × 17 columns



```
In [44]: df3=df3.drop(["date"],axis=1)
```

```
In [45]: sns.heatmap(df3.corr())
```

```
Out[45]: <AxesSubplot:>
```



```
In [46]: x=df3.drop(["TCH"],axis=1)
y=df3["TCH"]
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

## Linear


```
In [47]: li=LinearRegression()
li.fit(x_train,y_train)
```

```
Out[47]: LinearRegression()
```

```
In [ ]:
```

```
In [48]: prediction=li.predict(x_test)
plt.scatter(y_test,prediction)
```

```
Out[48]: <matplotlib.collections.PathCollection at 0x25551c2ebb0>
```



```
In [49]: lis=li.score(x_test,y_test)
```

```
In [50]: df3["TCH"].value_counts()
```

```
Out[50]: 1.34    740
         1.33    714
         1.35    708
         1.37    688
         1.36    679
         ...
         2.95     1
         3.65     1
         3.59     1
         2.58     1
         3.86     1
         Name: TCH, Length: 191, dtype: int64
```

```
In [51]: df3.loc[df3["TCH"]<1.40,"TCH"]=1
         df3.loc[df3["TCH"]>1.40,"TCH"]=2
         df3["TCH"].value_counts()
```

```
Out[51]: 1.0    11861
         2.0     7536
         Name: TCH, dtype: int64
```

```
In [ ]:
```


## Lasso

```
In [52]: la=Lasso(alpha=5)
         la.fit(x_train,y_train)
```

```
Out[52]: Lasso(alpha=5)
```

```
In [53]: prediction1=la.predict(x_test)
         plt.scatter(y_test,prediction1)
```

```
Out[53]: <matplotlib.collections.PathCollection at 0x25551d39cd0>
```



```
In [54]: las=la.score(x_test,y_test)
```


## Ridge

```
In [55]: rr=Ridge(alpha=1)  
rr.fit(x_train,y_train)
```

```
Out[55]: Ridge(alpha=1)
```

```
In [56]: prediction2=rr.predict(x_test)  
plt.scatter(y_test,prediction2)
```

```
Out[56]: <matplotlib.collections.PathCollection at 0x25551da22e0>
```



```
In [57]: rrs=rr.score(x_test,y_test)
```


## ElasticNet

```
In [58]: en=ElasticNet()  
en.fit(x_train,y_train)
```

```
Out[58]: ElasticNet()
```

```
In [59]: prediction2=rr.predict(x_test)
plt.scatter(y_test,prediction2)
```

```
Out[59]: <matplotlib.collections.PathCollection at 0x25551df9df0>
```



```
In [60]: ens=en.score(x_test,y_test)
```

```
In [61]: print(rr.score(x_test,y_test))
rr.score(x_train,y_train)
```

```
0.5759339018894761
```

```
Out[61]: 0.5960803646994859
```

## Logistic

```
In [62]: g={"TCH":{1.0:"Low",2.0:"High"}}
df3=df3.replace(g)
df3["TCH"].value_counts()
```

```
Out[62]: Low      11861
High       7536
Name: TCH, dtype: int64
```

```
In [63]: x=df3.drop(["TCH"],axis=1)
y=df3["TCH"]
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

```
In [64]: lo=LogisticRegression()
lo.fit(x_train,y_train)
```

```
Out[64]: LogisticRegression()
```



```
In [65]: prediction3=lo.predict(x_test)
plt.scatter(y_test,prediction3)
```

```
Out[65]: <matplotlib.collections.PathCollection at 0x25551710670>
```



```
In [66]: los=lo.score(x_test,y_test)
```

## Random Forest

```
In [67]: from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV
```

```
In [68]: g1={"TCH":{"Low":1.0,"High":2.0}}
df3=df3.replace(g1)
```

```
In [69]: x=df3.drop(["TCH"],axis=1)
y=df3["TCH"]
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

```
In [70]: rfc=RandomForestClassifier()
rfc.fit(x_train,y_train)
```

```
Out[70]: RandomForestClassifier()
```

```
In [71]: parameter={
    'max_depth':[1,2,4,5,6],
    'min_samples_leaf':[5,10,15,20,25],
    'n_estimators':[10,20,30,40,50]
}
```

```
In [72]: grid_search=GridSearchCV(estimator=rfc,param_grid=parameter,cv=2,scoring="accuracy")
grid_search.fit(x_train,y_train)
```

```
Out[72]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
    param_grid={'max_depth': [1, 2, 4, 5, 6],
    'min_samples_leaf': [5, 10, 15, 20, 25],
    'n_estimators': [10, 20, 30, 40, 50]},
    scoring='accuracy')
```

```
In [73]: rfcs=grid_search.best_score_
```

```
In [74]: rfc_best=grid_search.best_estimator_
```

```
In [75]: from sklearn.tree import plot_tree

plt.figure(figsize=(80,40))
plot_tree(rfc_best.estimators_[5],feature_names=x.columns,class_names=['Yes','No'],fillc
347\nvalue = [559, 36]\nnclass = Yes'),
  Text(192.41379310344828, 155.3142857142857, 'gini = 0.08\nsamples = 340\nvalue = [5
54, 24]\nnclass = Yes'),
  Text(269.3793103448276, 155.3142857142857, 'gini = 0.415\nsamples = 7\nvalue = [5,
12]\nnclass = No'),
  Text(461.79310344827593, 776.5714285714287, 'TOL <= 7.055\ngini = 0.236\nsamples =
1265\nvalue = [1709, 270]\nnclass = Yes'),
  Text(384.82758620689657, 465.9428571428573, 'PM25 <= 15.78\ngini = 0.202\nsamples =
1128\nvalue = [1559, 201]\nnclass = Yes'),
  Text(346.3448275862069, 155.3142857142857, 'gini = 0.111\nsamples = 777\nvalue = [1
129, 71]\nnclass = Yes'),
  Text(423.3103448275862, 155.3142857142857, 'gini = 0.357\nsamples = 351\nvalue = [4
30, 130]\nnclass = Yes'),
  Text(538.7586206896552, 465.9428571428573, 'SO_2 <= 8.61\ngini = 0.432\nsamples = 1
37\nvalue = [150, 69]\nnclass = Yes'),
  Text(500.2758620689656, 155.3142857142857, 'gini = 0.488\nsamples = 83\nvalue = [7
8, 57]\nnclass = Yes'),
  Text(577.2413793103449, 155.3142857142857, 'gini = 0.245\nsamples = 54\nvalue = [7
2, 12]\nnclass = Yes'),
  Text(923.5862068965519, 1087.2, 'BEN <= 2.135\ngini = 0.347\nsamples = 1327\nvalue
```

```
In [76]: print("Linear:",lis)
print("Lasso:",las)
print("Ridge:",rrs)
print("ElasticNet:",ens)
print("Logistic:",los)
print("Random Forest:",rfcs)
```

```
Linear: 0.5758911744718893
Lasso: 0.49886828067074285
Ridge: 0.5759339018894761
ElasticNet: 0.5250372976356534
Logistic: 0.6163230240549828
Random Forest: 0.8953377842749368
```

## Best model is Random Forest

```
In [ ]:
```

