

```
In [1]: import pandas as pd
import numpy as np
from matplotlib import pyplot as plt
import seaborn as sns
from sklearn.linear_model import LinearRegression, LogisticRegression, Lasso, Ridge, ElasticNet
from sklearn.model_selection import train_test_split
```

```
In [2]: df=pd.read_csv("C:/Users/user/Downloads/FP1_air/csvs_per_year/csvs_per_year/madrid_2007_08.csv")
df
```

Out[2]:

	date	BEN	CO	EBE	MXV	NMHC	NO_2	NOx	OXY	O_3	PM10	PM2.5
0	2007-12-01 01:00:00	NaN	2.86	NaN	NaN	NaN	282.200012	1054.000000	NaN	4.030000	156.199997	97.400000
1	2007-12-01 01:00:00	NaN	1.82	NaN	NaN	NaN	86.419998	354.600006	NaN	3.260000	80.809998	NaN
2	2007-12-01 01:00:00	NaN	1.47	NaN	NaN	NaN	94.639999	319.000000	NaN	5.310000	53.099998	NaN
3	2007-12-01 01:00:00	NaN	1.64	NaN	NaN	NaN	127.900002	476.700012	NaN	4.500000	105.300003	NaN
4	2007-12-01 01:00:00	4.64	1.86	4.26	7.98	0.57	145.100006	573.900024	3.49	52.689999	106.500000	15.900000
...
225115	2007-03-01 00:00:00	0.30	0.45	1.00	0.30	0.26	8.690000	11.690000	1.00	42.209999	6.760000	5.700000
225116	2007-03-01 00:00:00	NaN	0.16	NaN	NaN	NaN	46.820000	51.480000	NaN	22.150000	5.700000	NaN
225117	2007-03-01 00:00:00	0.24	NaN	0.20	NaN	0.09	51.259998	66.809998	NaN	18.540001	13.010000	6.900000
225118	2007-03-01 00:00:00	0.11	NaN	1.00	NaN	0.05	24.240000	36.930000	NaN	NaN	6.610000	NaN
225119	2007-03-01 00:00:00	0.53	0.40	1.00	1.70	0.12	32.360001	47.860001	1.37	24.150000	10.260000	7.000000

225120 rows × 17 columns

```
In [3]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 225120 entries, 0 to 225119
Data columns (total 17 columns):
 #   Column      Non-Null Count  Dtype  
---  -
 0   date        225120 non-null object  
 1   BEN         68885 non-null  float64
 2   CO          206748 non-null float64
 3   EBE         68883 non-null  float64
 4   MXY         26061 non-null  float64
 5   NMHC        86883 non-null  float64
 6   NO_2        223985 non-null float64
 7   NOx         223972 non-null float64
 8   OXY         26062 non-null  float64
 9   O_3         211850 non-null float64
10  PM10        222588 non-null float64
11  PM25        68870 non-null  float64
12  PXY         26062 non-null  float64
13  SO_2        224372 non-null float64
14  TCH         87026 non-null  float64
15  TOL         68845 non-null  float64
16  station     225120 non-null int64  
dtypes: float64(15), int64(1), object(1)
memory usage: 29.2+ MB
```

```
In [4]: df1=df.dropna()  
df1
```

Out[4]:


	date	BEN	CO	EBE	MXV	NMHC	NO_2	NOx	OXY	O_3	PM10	PM2.5
4	2007-12-01 01:00:00	4.64	1.86	4.26	7.98	0.57	145.100006	573.900024	3.49	52.689999	106.500000	15.900000
21	2007-12-01 01:00:00	1.98	0.31	2.56	6.06	0.35	76.059998	208.899994	1.70	1.000000	37.799999	25.500000
25	2007-12-01 01:00:00	2.82	1.42	3.15	7.02	0.49	123.099998	402.399994	2.60	7.160000	70.809998	37.000000
30	2007-12-01 02:00:00	4.65	1.89	4.41	8.21	0.65	151.000000	622.700012	3.55	58.080002	117.099998	17.040000
47	2007-12-01 02:00:00	1.97	0.30	2.15	5.08	0.33	78.760002	189.800003	1.62	1.000000	34.740002	24.700000
...
225073	2007-02-28 23:00:00	2.12	0.47	2.51	4.99	0.05	43.560001	83.889999	2.57	13.090000	21.860001	9.300000
225094	2007-02-28 23:00:00	0.87	0.45	1.19	2.66	0.13	40.000000	61.959999	1.79	20.440001	15.070000	9.200000
225098	2007-03-01 00:00:00	0.95	0.41	1.55	3.11	0.05	36.090000	63.349998	1.74	17.160000	9.210000	5.100000
225115	2007-03-01 00:00:00	0.30	0.45	1.00	0.30	0.26	8.690000	11.690000	1.00	42.209999	6.760000	5.140000
225119	2007-03-01 00:00:00	0.53	0.40	1.00	1.70	0.12	32.360001	47.860001	1.37	24.150000	10.260000	7.000000

25443 rows × 17 columns

```
In [5]: df1=df1.drop(["date"],axis=1)
```


```
In [6]: sns.heatmap(df1.corr())
```

```
Out[6]: <AxesSubplot:>
```



```
In [7]: plt.plot(df1["EBE"],df1["PXY"],"o")
```

```
Out[7]: [<matplotlib.lines.Line2D at 0x2118cb47160>]
```



```
In [8]: data=df[["EBE","PXY"]]
```

```
In [9]: # sns.stripplot(x=df["EBE"],y=df["PXY"],jitter=True,marker='o',color='blue')
```

```
In [10]: x=df1.drop(["EBE"],axis=1)
y=df1["EBE"]
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```


Linear

```
In [11]: li=LinearRegression()
li.fit(x_train,y_train)
```

```
Out[11]: LinearRegression()
```

```
In [12]: prediction=li.predict(x_test)
plt.scatter(y_test,prediction)
```

```
Out[12]: <matplotlib.collections.PathCollection at 0x2118d9e5bb0>
```



```
In [13]: lis=li.score(x_test,y_test)
```

```
In [14]: df1["TCH"].value_counts()
```

```
Out[14]: 1.34    1130
         1.33    1067
         1.35    1037
         1.36    1002
         1.32     991
         ...
         4.07      1
         2.71      1
         0.40      1
         0.38      1
         3.32      1
Name: TCH, Length: 250, dtype: int64
```

```
In [15]: df1.loc[df1["TCH"]<1.40,"TCH"]=1
df1.loc[df1["TCH"]>1.40,"TCH"]=2
df1["TCH"].value_counts()
```

```
Out[15]: 1.0    14025
         2.0    11418
Name: TCH, dtype: int64
```


Lasso

```
In [16]: la=Lasso(alpha=5)
la.fit(x_train,y_train)
```

```
Out[16]: Lasso(alpha=5)
```

```
In [17]: prediction1=la.predict(x_test)
plt.scatter(y_test,prediction1)
```

```
Out[17]: <matplotlib.collections.PathCollection at 0x2118da53880>
```



```
In [18]: las=la.score(x_test,y_test)
```


Ridge

```
In [19]: rr=Ridge(alpha=1)
rr.fit(x_train,y_train)
```

```
Out[19]: Ridge(alpha=1)
```

```
In [20]: prediction2=rr.predict(x_test)
plt.scatter(y_test,prediction2)
```

```
Out[20]: <matplotlib.collections.PathCollection at 0x2118cb1d6a0>
```



```
In [21]: rrs=rr.score(x_test,y_test)
```


ElasticNet

```
In [22]: en=ElasticNet()
en.fit(x_train,y_train)
```

```
Out[22]: ElasticNet()
```

```
In [23]: prediction2=rr.predict(x_test)
plt.scatter(y_test,prediction2)
```

```
Out[23]: <matplotlib.collections.PathCollection at 0x2118dade40>
```



```
In [24]: ens=en.score(x_test,y_test)
```

```
In [25]: print(rr.score(x_test,y_test))
rr.score(x_train,y_train)
```

```
0.9130065705546464
```

```
Out[25]: 0.8598350807252741
```

Logistic

```
In [26]: g={"TCH":{1.0:"Low",2.0:"High"}}
df1=df1.replace(g)
df1["TCH"].value_counts()
```

```
Out[26]: Low      14025
High      11418
Name: TCH, dtype: int64
```

```
In [27]: x=df1.drop(["TCH"],axis=1)
y=df1["TCH"]
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

```
In [28]: lo=LogisticRegression()
lo.fit(x_train,y_train)
```

```
Out[28]: LogisticRegression()
```

```
In [29]: prediction3=lo.predict(x_test)
plt.scatter(y_test,prediction3)
```

```
Out[29]: <matplotlib.collections.PathCollection at 0x2118db28d00>
```



```
In [30]: los=lo.score(x_test,y_test)
```

Random Forest

```
In [31]: from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV
```

```
In [32]: g1={"TCH":{"Low":1.0,"High":2.0}}
df1=df1.replace(g1)
```

```
In [33]: x=df1.drop(["TCH"],axis=1)
y=df1["TCH"]
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

```
In [34]: rfc=RandomForestClassifier()
rfc.fit(x_train,y_train)
```

```
Out[34]: RandomForestClassifier()
```

```
In [35]: parameter={
    'max_depth':[1,2,4,5,6],
    'min_samples_leaf':[5,10,15,20,25],
    'n_estimators':[10,20,30,40,50]
}
```

```
In [36]: grid_search=GridSearchCV(estimator=rfc,param_grid=parameter,cv=2,scoring="accuracy")
grid_search.fit(x_train,y_train)
```

```
Out[36]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
    param_grid={'max_depth': [1, 2, 4, 5, 6],
    'min_samples_leaf': [5, 10, 15, 20, 25],
    'n_estimators': [10, 20, 30, 40, 50]},
    scoring='accuracy')
```



```
In [37]: rfcs=grid_search.best_score_
```

```
In [38]: rfc_best=grid_search.best_estimator_
```

```
In [39]: from sklearn.tree import plot_tree

plt.figure(figsize=(80,40))
plot_tree(rfc_best.estimators_[5],feature_names=x.columns,class_names=['Yes','No'],fillc=
    Text(1450.8000000000002, 155.3142857142857, 'gini = 0.142\nsamples = 123\nvalue =
[16, 192]\nnclass = No'),
    Text(1785.6000000000001, 1087.2, 'CO <= 0.815\ngini = 0.065\nsamples = 2286\nvalue
= [123, 3535]\nnclass = No'),
    Text(1636.8000000000002, 776.5714285714287, 'PM25 <= 15.605\ngini = 0.119\nsamples
= 1103\nvalue = [112, 1656]\nnclass = No'),
    Text(1562.4, 465.9428571428573, 'OXY <= 2.695\ngini = 0.235\nsamples = 176\nvalue =
[37, 235]\nnclass = No'),
    Text(1525.2, 155.3142857142857, 'gini = 0.165\nsamples = 150\nvalue = [21, 211]\ncl
ass = No'),
    Text(1599.6000000000001, 155.3142857142857, 'gini = 0.48\nsamples = 26\nvalue = [1
6, 24]\nnclass = No'),
    Text(1711.2, 465.9428571428573, 'EBE <= 0.635\ngini = 0.095\nsamples = 927\nvalue =
[75, 1421]\nnclass = No'),
    Text(1674.0000000000002, 155.3142857142857, 'gini = 0.394\nsamples = 28\nvalue = [1
0, 27]\nnclass = No'),
    Text(1748.4, 155.3142857142857, 'gini = 0.085\nsamples = 899\nvalue = [65, 1394]\nc
lass = No'),
    Text(1934.4, 776.5714285714287, 'TOL <= 6.18\ngini = 0.012\nsamples = 1183\nvalue =
[11, 1879]\nnclass = No')
    
```

```
In [40]: print("Linear:",lis)
print("Lasso:",las)
print("Ridge:",rrs)
print("ElasticNet:",ens)
print("Logistic:",los)
print("Random Forest:",rfcs)
```

```
Linear: 0.9130054463649523
Lasso: 0.5416068633176341
Ridge: 0.9130065705546464
ElasticNet: 0.8500784330712808
Logistic: 0.5553517620856806
Random Forest: 0.8690061763054464
```

Best Model is Ridge Regression

```
In [41]: df2=pd.read_csv("C:/Users/user/Downloads/FP1_air/csvs_per_year/csvs_per_year/madrid_2008")
df2
```

Out[41]:

	date	BEN	CO	EBE	MXV	NMHC	NO_2	NOx	OXY	O_3	PM10	PM25
0	2008-06-01 01:00:00	NaN	0.47	NaN	NaN	NaN	83.089996	120.699997	NaN	16.990000	16.889999	10.40
1	2008-06-01 01:00:00	NaN	0.59	NaN	NaN	NaN	94.820000	130.399994	NaN	17.469999	19.040001	NaN
2	2008-06-01 01:00:00	NaN	0.55	NaN	NaN	NaN	75.919998	104.599998	NaN	13.470000	20.270000	NaN
3	2008-06-01 01:00:00	NaN	0.36	NaN	NaN	NaN	61.029999	66.559998	NaN	23.110001	10.850000	NaN
4	2008-06-01 01:00:00	1.68	0.80	1.70	3.01	0.30	105.199997	214.899994	1.61	12.120000	37.160000	21.90
...
226387	2008-11-01 00:00:00	0.48	0.30	0.57	1.00	0.31	13.050000	14.160000	0.91	57.400002	5.450000	5.15
226388	2008-11-01 00:00:00	NaN	0.30	NaN	NaN	NaN	41.880001	48.500000	NaN	35.830002	15.020000	NaN
226389	2008-11-01 00:00:00	0.25	NaN	0.56	NaN	0.11	83.610001	102.199997	NaN	14.130000	17.540001	13.91
226390	2008-11-01 00:00:00	0.54	NaN	2.70	NaN	0.18	70.639999	81.860001	NaN	NaN	11.910000	NaN
226391	2008-11-01 00:00:00	0.75	0.36	1.20	2.75	0.16	58.240002	74.239998	1.64	31.910000	12.690000	11.42

226392 rows × 17 columns



```
In [42]: df2.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 226392 entries, 0 to 226391
Data columns (total 17 columns):
 #   Column      Non-Null Count  Dtype  
---  -
 0   date        226392 non-null object  
 1   BEN         67047 non-null float64
 2   CO          208109 non-null float64
 3   EBE         67044 non-null float64
 4   MXY         25867 non-null float64
 5   NMHC        85079 non-null float64
 6   NO_2        225315 non-null float64
 7   NOx         225311 non-null float64
 8   OXY         25878 non-null float64
 9   O_3         215716 non-null float64
10  PM10        220179 non-null float64
11  PM25        67833 non-null float64
12  PXY         25877 non-null float64
13  SO_2        225405 non-null float64
14  TCH         85107 non-null float64
15  TOL         66940 non-null float64
16  station     226392 non-null int64  
dtypes: float64(15), int64(1), object(1)
memory usage: 29.4+ MB
```

In [43]:

df3=df2.dropna()
df3

Out[43]:

	date	BEN	CO	EBE	MXV	NMHC	NO_2	NOx	OXY	O_3	PM10	P
4	2008-06-01 01:00:00	1.68	0.80	1.70	3.01	0.30	105.199997	214.899994	1.61	12.120000	37.160000	21.900
21	2008-06-01 01:00:00	0.32	0.37	1.00	0.39	0.33	21.580000	22.180000	1.00	35.770000	7.900000	6.140
25	2008-06-01 01:00:00	0.73	0.39	1.04	1.70	0.18	64.839996	86.709999	1.31	23.379999	14.760000	9.840
30	2008-06-01 02:00:00	1.95	0.51	1.98	3.77	0.24	79.750000	143.399994	2.03	18.090000	31.139999	18.410
47	2008-06-01 02:00:00	0.36	0.39	0.39	0.50	0.34	26.790001	27.389999	1.00	33.029999	7.620000	6.250
...
226362	2008-10-31 23:00:00	0.47	0.35	0.65	1.00	0.33	22.480000	25.020000	1.00	33.509998	10.200000	7.680
226366	2008-10-31 23:00:00	0.92	0.46	1.21	2.75	0.19	78.440002	106.199997	1.70	18.320000	14.140000	10.590
226371	2008-11-01 00:00:00	1.83	0.53	2.22	4.51	0.17	93.260002	158.399994	2.38	18.770000	20.750000	18.620
226387	2008-11-01 00:00:00	0.48	0.30	0.57	1.00	0.31	13.050000	14.160000	0.91	57.400002	5.450000	5.150
226391	2008-11-01 00:00:00	0.75	0.36	1.20	2.75	0.16	58.240002	74.239998	1.64	31.910000	12.690000	11.420


25631 rows × 17 columns

In [44]:

df3=df3.drop(["date"],axis=1)

```
In [45]: sns.heatmap(df3.corr())
```

```
Out[45]: <AxesSubplot:>
```



```
In [46]: x=df3.drop(["TCH"],axis=1)
y=df3["TCH"]
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

Linear


```
In [47]: li=LinearRegression()
li.fit(x_train,y_train)
```

```
Out[47]: LinearRegression()
```

```
In [ ]:
```

```
In [48]: prediction=li.predict(x_test)
plt.scatter(y_test,prediction)
```

```
Out[48]: <matplotlib.collections.PathCollection at 0x2118db65550>
```



```
In [49]: lis=li.score(x_test,y_test)
```

```
In [50]: df3["TCH"].value_counts()
```

```
Out[50]: 1.38    1274
         1.37    1246
         1.36    1243
         1.39    1242
         1.35    1209
         ...
         2.41     1
         2.95     1
         0.98     1
         2.64     1
         2.61     1
         Name: TCH, Length: 177, dtype: int64
```

```
In [51]: df3.loc[df3["TCH"]<1.40,"TCH"]=1
         df3.loc[df3["TCH"]>1.40,"TCH"]=2
         df3["TCH"].value_counts()
```

```
Out[51]: 2.0    12904
         1.0    12727
         Name: TCH, dtype: int64
```

```
In [ ]:
```


Lasso

```
In [52]: la=Lasso(alpha=5)
         la.fit(x_train,y_train)
```

```
Out[52]: Lasso(alpha=5)
```

```
In [53]: prediction1=la.predict(x_test)
         plt.scatter(y_test,prediction1)
```

```
Out[53]: <matplotlib.collections.PathCollection at 0x2118dbbc100>
```



```
In [54]: las=la.score(x_test,y_test)
```


Ridge

```
In [55]: rr=Ridge(alpha=1)  
rr.fit(x_train,y_train)
```

```
Out[55]: Ridge(alpha=1)
```

```
In [56]: prediction2=rr.predict(x_test)  
plt.scatter(y_test,prediction2)
```

```
Out[56]: <matplotlib.collections.PathCollection at 0x2118dc096d0>
```



```
In [57]: rrs=rr.score(x_test,y_test)
```


ElasticNet

```
In [58]: en=ElasticNet()  
en.fit(x_train,y_train)
```

```
Out[58]: ElasticNet()
```

```
In [59]: prediction2=rr.predict(x_test)
plt.scatter(y_test,prediction2)
```

```
Out[59]: <matplotlib.collections.PathCollection at 0x2118dc5cca0>
```



```
In [60]: ens=en.score(x_test,y_test)
```

```
In [61]: print(rr.score(x_test,y_test))
rr.score(x_train,y_train)
```

```
0.6614787517671646
```

```
Out[61]: 0.6579137146941019
```

Logistic

```
In [62]: g={"TCH":{1.0:"Low",2.0:"High"}}
df3=df3.replace(g)
df3["TCH"].value_counts()
```

```
Out[62]: High    12904
Low      12727
Name: TCH, dtype: int64
```

```
In [63]: x=df3.drop(["TCH"],axis=1)
y=df3["TCH"]
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

```
In [64]: lo=LogisticRegression()
lo.fit(x_train,y_train)
```

```
Out[64]: LogisticRegression()
```



```
In [65]: prediction3=lo.predict(x_test)
plt.scatter(y_test,prediction3)
```

```
Out[65]: <matplotlib.collections.PathCollection at 0x2118d472880>
```



```
In [66]: los=lo.score(x_test,y_test)
```

Random Forest

```
In [67]: from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV
```

```
In [68]: g1={"TCH":{"Low":1.0,"High":2.0}}
df3=df3.replace(g1)
```

```
In [69]: x=df3.drop(["TCH"],axis=1)
y=df3["TCH"]
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

```
In [70]: rfc=RandomForestClassifier()
rfc.fit(x_train,y_train)
```

```
Out[70]: RandomForestClassifier()
```

```
In [71]: parameter={
    'max_depth':[1,2,4,5,6],
    'min_samples_leaf':[5,10,15,20,25],
    'n_estimators':[10,20,30,40,50]
}
```

```
In [72]: grid_search=GridSearchCV(estimator=rfc,param_grid=parameter,cv=2,scoring="accuracy")
grid_search.fit(x_train,y_train)
```

```
Out[72]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
    param_grid={'max_depth': [1, 2, 4, 5, 6],
    'min_samples_leaf': [5, 10, 15, 20, 25],
    'n_estimators': [10, 20, 30, 40, 50]},
    scoring='accuracy')
```

```
In [73]: rfcs=grid_search.best_score_
```

```
In [74]: rfc_best=grid_search.best_estimator_
```

```
In [75]: from sklearn.tree import plot_tree
```

```
plt.figure(figsize=(80,40))
plot_tree(rfc_best.estimators_[5],feature_names=x.columns,class_names=['Yes','No'],fill
```

```
Out[75]: [Text(2217.053571428571, 2019.0857142857144, 'NOx <= 95.465\ngini = 0.5\nsamples = 1
1390\nvalue = [8921, 9020]\nclass = No'),
Text(1101.0535714285713, 1708.457142857143, 'NO_2 <= 45.97\ngini = 0.411\nsamples =
7185\nvalue = [8055, 3272]\nclass = Yes'),
Text(548.0357142857142, 1397.8285714285716, 'station <= 28079015.0\ngini = 0.364\ns
amples = 5401\nvalue = [6447, 2024]\nclass = Yes'),
Text(298.9285714285714, 1087.2, 'NMHC <= 0.255\ngini = 0.162\nsamples = 735\nvalue
= [1047, 102]\nclass = Yes'),
Text(159.42857142857142, 776.5714285714287, 'MXY <= 1.275\ngini = 0.107\nsamples =
703\nvalue = [1043, 63]\nclass = Yes'),
Text(79.71428571428571, 465.9428571428573, 'CO <= 0.425\ngini = 0.035\nsamples = 43
2\nvalue = [652, 12]\nclass = Yes'),
Text(39.857142857142854, 155.3142857142857, 'gini = 0.024\nsamples = 426\nvalue =
[648, 8]\nclass = Yes'),
Text(119.57142857142856, 155.3142857142857, 'gini = 0.5\nsamples = 6\nvalue = [4,
4]\nclass = Yes'),
Text(239.1428571428571, 465.9428571428573, 'NO_2 <= 36.62\ngini = 0.204\nsamples =
271\nvalue = [391, 51]\nclass = Yes'),
Text(199.28571428571428, 155.3142857142857, 'gini = 0.105\nsamples = 130\nvalue =
[302, 12]\nclass = Yes')]
```

```
In [76]: print("Linear:",lis)
print("Lasso:",las)
print("Ridge:",rrs)
print("ElasticNet:",ens)
print("Logistic:",los)
print("Random Forest:",rfcs)
```

```
Linear: 0.6614554468063345
Lasso: 0.4603230758526199
Ridge: 0.6614787517671646
ElasticNet: 0.5801002117823061
Logistic: 0.5009102730819246
Random Forest: 0.8317820819146347
```

Best model is Random Forest

```
In [ ]:
```

