

```
In [1]: import pandas as pd
import numpy as np
from matplotlib import pyplot as plt
import seaborn as sns
from sklearn.linear_model import LinearRegression, LogisticRegression, Lasso, Ridge, ElasticNet
from sklearn.model_selection import train_test_split
```

```
In [2]: df=pd.read_csv("C:/Users/user/Downloads/FP1_air/csvs_per_year/csvs_per_year/madrid_2017")
df
```

Out[2]:

	date	BEN	CH4	CO	EBE	NMHC	NO	NO_2	NOx	O_3	PM10	PM25	SO_2	TCH	TOL	
0	2017-06-01 01:00:00	NaN	NaN	0.3	NaN	NaN	4.0	38.0	NaN	NaN	NaN	NaN	5.0	NaN	NaN	28
1	2017-06-01 01:00:00	0.6	NaN	0.3	0.4	0.08	3.0	39.0	NaN	71.0	22.0	9.0	7.0	1.4	2.9	28
2	2017-06-01 01:00:00	0.2	NaN	NaN	0.1	NaN	1.0	14.0	NaN	NaN	NaN	NaN	NaN	NaN	0.9	28
3	2017-06-01 01:00:00	NaN	NaN	0.2	NaN	NaN	1.0	9.0	NaN	91.0	NaN	NaN	NaN	NaN	NaN	28
4	2017-06-01 01:00:00	NaN	NaN	NaN	NaN	NaN	1.0	19.0	NaN	69.0	NaN	NaN	2.0	NaN	NaN	28
...
210115	2017-08-01 00:00:00	NaN	NaN	0.2	NaN	NaN	1.0	27.0	NaN	65.0	NaN	NaN	NaN	NaN	NaN	28
210116	2017-08-01 00:00:00	NaN	NaN	0.2	NaN	NaN	1.0	14.0	NaN	NaN	73.0	NaN	7.0	NaN	NaN	28
210117	2017-08-01 00:00:00	NaN	NaN	NaN	NaN	NaN	1.0	4.0	NaN	83.0	NaN	NaN	NaN	NaN	NaN	28
210118	2017-08-01 00:00:00	NaN	NaN	NaN	NaN	NaN	1.0	11.0	NaN	78.0	NaN	NaN	NaN	NaN	NaN	28
210119	2017-08-01 00:00:00	NaN	NaN	NaN	NaN	NaN	1.0	14.0	NaN	77.0	60.0	NaN	NaN	NaN	NaN	28

210120 rows × 16 columns



```
In [3]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 210120 entries, 0 to 210119
Data columns (total 16 columns):
 #   Column      Non-Null Count  Dtype  
---  -
 0   date        210120 non-null object  
 1   BEN         50201 non-null float64
 2   CH4         6410 non-null  float64
 3   CO          87001 non-null float64
 4   EBE         49973 non-null float64
 5   NMHC        25472 non-null float64
 6   NO          209065 non-null float64
 7   NO_2        209065 non-null float64
 8   NOx         52818 non-null float64
 9   O_3         121398 non-null float64
10  PM10        104141 non-null float64
11  PM25        52023 non-null float64
12  SO_2        86803 non-null float64
13  TCH         25472 non-null float64
14  TOL         50117 non-null float64
15  station     210120 non-null int64  
dtypes: float64(14), int64(1), object(1)
memory usage: 25.6+ MB
```

```
In [4]: df1=df.dropna()
df1
```

Out[4]:

	date	BEN	CH4	CO	EBE	NMHC	NO	NO_2	NOx	O_3	PM10	PM25	SO_2	TCH	TOL	s
87457	2017-10-01 01:00:00	0.6	1.22	0.3	0.4	0.09	4.0	54.0	60.0	43.0	12.0	9.0	13.0	1.31	2.3	280
87462	2017-10-01 01:00:00	0.2	1.18	0.2	0.1	0.09	1.0	26.0	28.0	42.0	14.0	6.0	3.0	1.27	1.1	280
87481	2017-10-01 02:00:00	0.4	1.22	0.2	0.2	0.06	2.0	32.0	36.0	53.0	14.0	10.0	13.0	1.28	1.3	280
87486	2017-10-01 02:00:00	0.2	1.19	0.2	0.1	0.07	1.0	15.0	17.0	51.0	18.0	8.0	3.0	1.26	0.8	280
87505	2017-10-01 03:00:00	0.3	1.23	0.2	0.2	0.06	2.0	27.0	29.0	57.0	15.0	10.0	13.0	1.29	1.0	280
...
158238	2017-12-31 22:00:00	0.3	1.11	0.2	0.1	0.03	1.0	8.0	9.0	73.0	3.0	1.0	3.0	1.14	0.2	280
158257	2017-12-31 23:00:00	0.6	1.38	0.3	0.1	0.03	6.0	42.0	51.0	47.0	7.0	4.0	3.0	1.41	0.9	280
158262	2017-12-31 23:00:00	0.3	1.11	0.2	0.1	0.03	1.0	6.0	8.0	72.0	6.0	3.0	3.0	1.14	0.2	280
158281	2018-01-01 00:00:00	0.5	1.38	0.2	0.1	0.02	2.0	20.0	23.0	69.0	4.0	2.0	3.0	1.39	0.6	280
158286	2018-01-01 00:00:00	0.3	1.11	0.2	0.1	0.03	1.0	1.0	3.0	83.0	8.0	5.0	3.0	1.14	0.2	280


4127 rows × 16 columns



```
In [5]: df1=df1.drop(["date"],axis=1)
```


```
In [6]: sns.heatmap(df1.corr())
```

```
Out[6]: <AxesSubplot:>
```



```
In [7]: plt.plot(df1["EBE"],df1["PM25"],"o")
```

```
Out[7]: [<matplotlib.lines.Line2D at 0x25f5f6d7fd0>]
```



```
In [8]: x=df1.drop(["EBE"],axis=1)
y=df1["EBE"]
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```


Linear

```
In [9]: li=LinearRegression()
li.fit(x_train,y_train)
```

```
Out[9]: LinearRegression()
```

```
In [10]: prediction=li.predict(x_test)
plt.scatter(y_test,prediction)
```

```
Out[10]: <matplotlib.collections.PathCollection at 0x25f5f7a3c70>
```



```
In [11]: lis=li.score(x_test,y_test)
```

```
In [12]: df1["TCH"].value_counts()
```

```
Out[12]: 1.24    124
         1.36    118
         1.26    112
         1.25    110
         1.41    107
         ...
         3.23     1
         2.47     1
         2.35     1
         2.61     1
         2.94     1
Name: TCH, Length: 164, dtype: int64
```

```
In [13]: df1.loc[df1["TCH"]<1.40,"TCH"]=1
df1.loc[df1["TCH"]>1.40,"TCH"]=2
df1["TCH"].value_counts()
```

```
Out[13]: 1.0    2428
         2.0    1699
Name: TCH, dtype: int64
```


Lasso

```
In [14]: la=Lasso(alpha=5)
la.fit(x_train,y_train)
```

```
Out[14]: Lasso(alpha=5)
```

```
In [15]: prediction1=la.predict(x_test)
plt.scatter(y_test,prediction1)
```

```
Out[15]: <matplotlib.collections.PathCollection at 0x25f5f80f820>
```



```
In [16]: las=la.score(x_test,y_test)
```


Ridge

```
In [17]: rr=Ridge(alpha=1)
rr.fit(x_train,y_train)
```

```
Out[17]: Ridge(alpha=1)
```

```
In [18]: prediction2=rr.predict(x_test)
plt.scatter(y_test,prediction2)
```

```
Out[18]: <matplotlib.collections.PathCollection at 0x25f5f64b2e0>
```



```
In [19]: rrs=rr.score(x_test,y_test)
```


ElasticNet

```
In [20]: en=ElasticNet()
en.fit(x_train,y_train)
```

```
Out[20]: ElasticNet()
```

```
In [21]: prediction2=rr.predict(x_test)
plt.scatter(y_test,prediction2)
```

```
Out[21]: <matplotlib.collections.PathCollection at 0x25f60140640>
```



```
In [22]: ens=en.score(x_test,y_test)
```

```
In [23]: print(rr.score(x_test,y_test))
rr.score(x_train,y_train)
```

```
0.8446472469016151
```

```
Out[23]: 0.8976558488919304
```

Logistic

```
In [24]: g={"TCH":{1.0:"Low",2.0:"High"}}
df1=df1.replace(g)
df1["TCH"].value_counts()
```

```
Out[24]: Low      2428
High      1699
Name: TCH, dtype: int64
```

```
In [25]: x=df1.drop(["TCH"],axis=1)
y=df1["TCH"]
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

```
In [26]: lo=LogisticRegression()
lo.fit(x_train,y_train)
```

```
Out[26]: LogisticRegression()
```

```
In [27]: prediction3=lo.predict(x_test)
plt.scatter(y_test,prediction3)
```

```
Out[27]: <matplotlib.collections.PathCollection at 0x25f601ad700>
```



```
In [28]: los=lo.score(x_test,y_test)
```

Random Forest

```
In [29]: from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV
```

```
In [30]: g1={"TCH":{"Low":1.0,"High":2.0}}
df1=df1.replace(g1)
```

```
In [31]: x=df1.drop(["TCH"],axis=1)
y=df1["TCH"]
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

```
In [32]: rfc=RandomForestClassifier()
rfc.fit(x_train,y_train)
```

```
Out[32]: RandomForestClassifier()
```

```
In [33]: parameter={
    'max_depth':[1,2,4,5,6],
    'min_samples_leaf':[5,10,15,20,25],
    'n_estimators':[10,20,30,40,50]
}
```

```
In [34]: grid_search=GridSearchCV(estimator=rfc,param_grid=parameter,cv=2,scoring="accuracy")
grid_search.fit(x_train,y_train)
```

```
Out[34]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
    param_grid={'max_depth': [1, 2, 4, 5, 6],
    'min_samples_leaf': [5, 10, 15, 20, 25],
    'n_estimators': [10, 20, 30, 40, 50]},
    scoring='accuracy')
```



```
In [35]: rfcs=grid_search.best_score_
```

```
In [36]: rfc_best=grid_search.best_estimator_
```

```
In [37]: from sklearn.tree import plot_tree

plt.figure(figsize=(80,40))
plot_tree(rfc_best.estimators_[5],feature_names=x.columns,class_names=['Yes','No'],fill
```

```
Out[37]: [Text(1962.6206896551726, 2019.0857142857144, 'BEN <= 0.65\ngini = 0.477\nsamples = 1820\nvalue = [1757, 1131]\nclass = Yes'),
Text(837.0000000000001, 1708.457142857143, 'NO_2 <= 24.5\ngini = 0.239\nsamples = 940\nvalue = [1312, 211]\nclass = Yes'),
Text(384.82758620689657, 1397.8285714285716, 'CH4 <= 1.315\ngini = 0.069\nsamples = 362\nvalue = [567, 21]\nclass = Yes'),
Text(230.89655172413796, 1087.2, 'CO <= 0.15\ngini = 0.004\nsamples = 342\nvalue = [555, 1]\nclass = Yes'),
Text(153.93103448275863, 776.5714285714287, 'CH4 <= 1.215\ngini = 0.021\nsamples = 60\nvalue = [93, 1]\nclass = Yes'),
Text(76.96551724137932, 465.9428571428573, 'gini = 0.0\nsamples = 55\nvalue = [89, 0]\nclass = Yes'),
Text(230.89655172413796, 465.9428571428573, 'gini = 0.32\nsamples = 5\nvalue = [4, 1]\nclass = Yes'),
Text(307.86206896551727, 776.5714285714287, 'gini = 0.0\nsamples = 282\nvalue = [462, 0]\nclass = Yes'),
Text(538.7586206896552, 1087.2, 'NMHC <= 0.025\ngini = 0.469\nsamples = 20\nvalue = [12, 20]\nclass = No'),
Text(461.79310344827593, 776.5714285714287, 'TOL <= 0.75\ngini = 0.245\nsamples = 10\nvalue = [12, 20]\nclass = Yes')]
```

```
In [38]: print("Linear:",lis)
print("Lasso:",las)
print("Ridge:",rrs)
print("ElasticNet:",ens)
print("Logistic:",los)
print("Random Forest:",rfcs)
```

```
Linear: 0.8441268043977549
Lasso: 0.6393010463991731
Ridge: 0.8446472469016151
ElasticNet: 0.7899757472729623
Logistic: 0.579499596448749
Random Forest: 0.9681440443213296
```

Best Model is Random Forest

```
In [39]: df2=pd.read_csv("C:/Users/user/Downloads/FP1_air/csvs_per_year/csvs_per_year/madrid_2018")
df2
```

Out[39]:

	date	BEN	CH4	CO	EBE	NMHC	NO	NO_2	NOx	O_3	PM10	PM25	SO_2	TCH	TOL
0	2018-03-01 01:00:00	NaN	NaN	0.3	NaN	NaN	1.0	29.0	31.0	NaN	NaN	NaN	2.0	NaN	NaN
1	2018-03-01 01:00:00	0.5	1.39	0.3	0.2	0.02	6.0	40.0	49.0	52.0	5.0	4.0	3.0	1.41	0.8
2	2018-03-01 01:00:00	0.4	NaN	NaN	0.2	NaN	4.0	41.0	47.0	NaN	NaN	NaN	NaN	NaN	1.1
3	2018-03-01 01:00:00	NaN	NaN	0.3	NaN	NaN	1.0	35.0	37.0	54.0	NaN	NaN	NaN	NaN	NaN
4	2018-03-01 01:00:00	NaN	NaN	NaN	NaN	NaN	1.0	27.0	29.0	49.0	NaN	NaN	3.0	NaN	NaN
...
69091	2018-02-01 00:00:00	NaN	NaN	0.5	NaN	NaN	66.0	91.0	192.0	1.0	35.0	22.0	NaN	NaN	NaN
69092	2018-02-01 00:00:00	NaN	NaN	0.7	NaN	NaN	87.0	107.0	241.0	NaN	29.0	NaN	15.0	NaN	NaN
69093	2018-02-01 00:00:00	NaN	NaN	NaN	NaN	NaN	28.0	48.0	91.0	2.0	NaN	NaN	NaN	NaN	NaN
69094	2018-02-01 00:00:00	NaN	NaN	NaN	NaN	NaN	141.0	103.0	320.0	2.0	NaN	NaN	NaN	NaN	NaN
69095	2018-02-01 00:00:00	NaN	NaN	NaN	NaN	NaN	69.0	96.0	202.0	3.0	26.0	NaN	NaN	NaN	NaN

69096 rows × 16 columns

```
In [40]: df2.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 69096 entries, 0 to 69095
Data columns (total 16 columns):
 #   Column      Non-Null Count  Dtype  
---  -
 0   date        69096 non-null  object  
 1   BEN         16950 non-null  float64 
 2   CH4         8440 non-null   float64 
 3   CO          28598 non-null  float64 
 4   EBE         16949 non-null  float64 
 5   NMHC        8440 non-null   float64 
 6   NO          68826 non-null  float64 
 7   NO_2        68826 non-null  float64 
 8   NOx         68826 non-null  float64 
 9   O_3         40049 non-null  float64 
10  PM10        36911 non-null  float64 
11  PM25        18912 non-null  float64 
12  SO_2        28586 non-null  float64 
13  TCH         8440 non-null   float64 
14  TOL         16950 non-null  float64 
15  station     69096 non-null  int64   
dtypes: float64(14), int64(1), object(1)
memory usage: 8.4+ MB
```

In [41]:

df3=df2.dropna()
df3

Out[41]:

	date	BEN	CH4	CO	EBE	NMHC	NO	NO_2	NOx	O_3	PM10	PM25	SO_2	TCH	TOL	
1	2018-03-01 01:00:00	0.5	1.39	0.3	0.2	0.02	6.0	40.0	49.0	52.0	5.0	4.0	3.0	1.41	0.8	28
6	2018-03-01 01:00:00	0.4	1.11	0.2	0.1	0.06	1.0	25.0	27.0	55.0	5.0	4.0	4.0	1.16	1.4	28
25	2018-03-01 02:00:00	0.4	1.42	0.2	0.1	0.01	4.0	26.0	32.0	64.0	4.0	4.0	3.0	1.44	0.7	28
30	2018-03-01 02:00:00	0.3	1.10	0.2	0.1	0.05	1.0	12.0	13.0	69.0	5.0	4.0	4.0	1.14	0.8	28
49	2018-03-01 03:00:00	0.3	1.41	0.2	0.1	0.01	3.0	16.0	20.0	68.0	3.0	2.0	3.0	1.42	0.4	28
...
69030	2018-01-31 22:00:00	1.8	1.21	0.7	1.7	0.19	151.0	129.0	361.0	1.0	45.0	26.0	11.0	1.40	11.9	28
69049	2018-01-31 23:00:00	3.1	1.87	1.2	2.0	0.35	296.0	162.0	615.0	3.0	39.0	23.0	8.0	2.22	12.5	28
69054	2018-01-31 23:00:00	1.6	1.17	0.6	1.4	0.15	127.0	106.0	301.0	1.0	43.0	25.0	8.0	1.32	10.3	28
69073	2018-02-01 00:00:00	3.2	1.53	1.0	2.1	0.19	125.0	117.0	309.0	3.0	37.0	24.0	6.0	1.72	13.0	28
69078	2018-02-01 00:00:00	1.3	1.14	0.4	0.8	0.10	54.0	73.0	155.0	1.0	27.0	16.0	5.0	1.24	6.8	28


4562 rows × 16 columns

In [42]:

df3=df3.drop(["date"],axis=1)

```
In [43]: sns.heatmap(df3.corr())
```

```
Out[43]: <AxesSubplot:>
```



```
In [44]: x=df3.drop(["TCH"],axis=1)
y=df3["TCH"]
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

Linear


```
In [45]: li=LinearRegression()
li.fit(x_train,y_train)
```

```
Out[45]: LinearRegression()
```

```
In [ ]:
```

```
In [46]: prediction=li.predict(x_test)
plt.scatter(y_test,prediction)
```

```
Out[46]: <matplotlib.collections.PathCollection at 0x25f605d5070>
```



```
In [47]: lis=li.score(x_test,y_test)
```

```
In [48]: df3["TCH"].value_counts()
```

```
Out[48]: 1.15    246
         1.43    232
         1.44    223
         1.14    210
         1.13    201
         ...
         2.35     1
         2.58     1
         2.73     1
         2.12     1
         1.96     1
         Name: TCH, Length: 143, dtype: int64
```

```
In [49]: df3.loc[df3["TCH"]<1.40,"TCH"]=1
         df3.loc[df3["TCH"]>1.40,"TCH"]=2
         df3["TCH"].value_counts()
```

```
Out[49]: 2.0    2477
         1.0    2085
         Name: TCH, dtype: int64
```

```
In [ ]:
```


Lasso

```
In [50]: la=Lasso(alpha=5)
         la.fit(x_train,y_train)
```

```
Out[50]: Lasso(alpha=5)
```

```
In [51]: prediction1=la.predict(x_test)
         plt.scatter(y_test,prediction1)
```

```
Out[51]: <matplotlib.collections.PathCollection at 0x25f60639190>
```



```
In [52]: las=la.score(x_test,y_test)
```


Ridge

```
In [53]: rr=Ridge(alpha=1)  
rr.fit(x_train,y_train)
```

```
Out[53]: Ridge(alpha=1)
```

```
In [54]: prediction2=rr.predict(x_test)  
plt.scatter(y_test,prediction2)
```

```
Out[54]: <matplotlib.collections.PathCollection at 0x25f6068d760>
```



```
In [55]: rrs=rr.score(x_test,y_test)
```


ElasticNet

```
In [56]: en=ElasticNet()  
en.fit(x_train,y_train)
```

```
Out[56]: ElasticNet()
```

```
In [57]: prediction2=rr.predict(x_test)
plt.scatter(y_test,prediction2)
```

```
Out[57]: <matplotlib.collections.PathCollection at 0x25f606f32b0>
```



```
In [58]: ens=en.score(x_test,y_test)
```

```
In [59]: print(rr.score(x_test,y_test))
rr.score(x_train,y_train)
```

```
0.9979283454674264
```

```
Out[59]: 0.998113461617783
```

Logistic

```
In [60]: g={"TCH":{1.0:"Low",2.0:"High"}}
df3=df3.replace(g)
df3["TCH"].value_counts()
```

```
Out[60]: High    2477
Low      2085
Name: TCH, dtype: int64
```

```
In [61]: x=df3.drop(["TCH"],axis=1)
y=df3["TCH"]
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```


```
In [62]: lo=LogisticRegression()
lo.fit(x_train,y_train)
```

```
Out[62]: LogisticRegression()
```



```
In [63]: prediction3=lo.predict(x_test)
plt.scatter(y_test,prediction3)
```

```
Out[63]: <matplotlib.collections.PathCollection at 0x25f6075c7f0>
```



```
In [64]: los=lo.score(x_test,y_test)
```

Random Forest

```
In [65]: from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV
```

```
In [66]: g1={"TCH":{"Low":1.0,"High":2.0}}
df3=df3.replace(g1)
```

```
In [67]: x=df3.drop(["TCH"],axis=1)
y=df3["TCH"]
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

```
In [68]: rfc=RandomForestClassifier()
rfc.fit(x_train,y_train)
```

```
Out[68]: RandomForestClassifier()
```

```
In [69]: parameter={
    'max_depth':[1,2,4,5,6],
    'min_samples_leaf':[5,10,15,20,25],
    'n_estimators':[10,20,30,40,50]
}
```

```
In [70]: grid_search=GridSearchCV(estimator=rfc,param_grid=parameter,cv=2,scoring="accuracy")
grid_search.fit(x_train,y_train)
```

```
Out[70]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
    param_grid={'max_depth': [1, 2, 4, 5, 6],
    'min_samples_leaf': [5, 10, 15, 20, 25],
    'n_estimators': [10, 20, 30, 40, 50]},
    scoring='accuracy')
```

```
In [71]: rfcs=grid_search.best_score_
```

```
In [72]: rfc_best=grid_search.best_estimator_
```

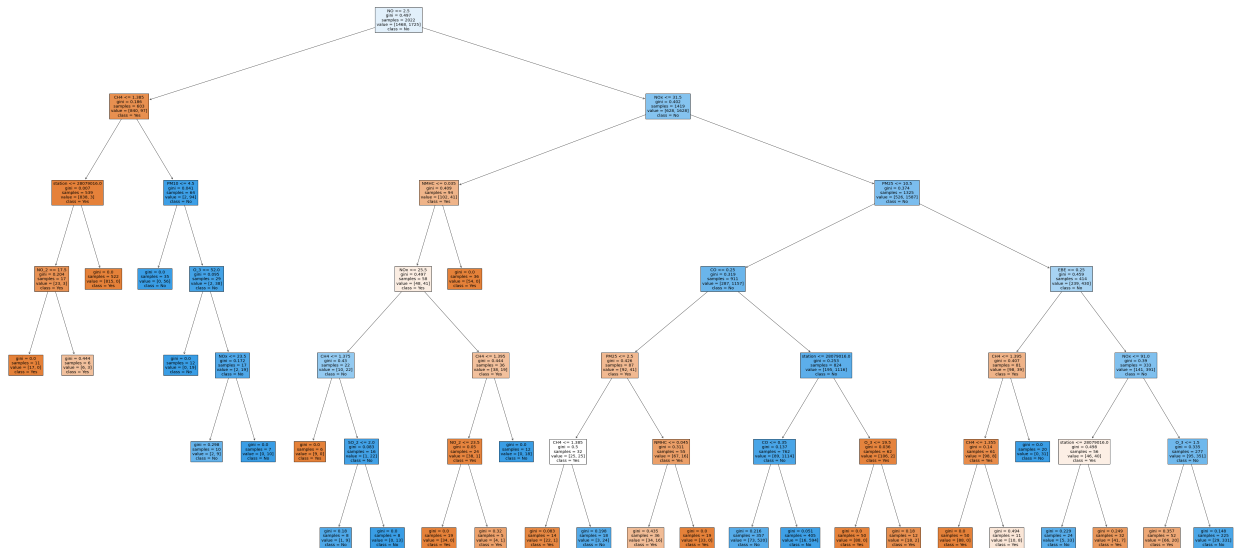
```
In [73]: from sklearn.tree import plot_tree  
  
plt.figure(figsize=(80,40))  
plot_tree(rfc_best.estimators_[5],feature_names=x.columns,class_names=['Yes','No'],fill,
```

```

Out[73]: [Text(1435.6875, 2019.0857142857144, 'NO <= 2.5\ngini = 0.497\nsamples = 2022\nvalue =
[1468, 1725]\nnclass = No'),
Text(465.0, 1708.457142857143, 'CH4 <= 1.385\ngini = 0.186\nsamples = 603\nvalue = [8
40, 97]\nnclass = Yes'),
Text(279.0, 1397.8285714285716, 'station <= 28079016.0\ngini = 0.007\nsamples = 539\n
value = [838, 3]\nnclass = Yes'),
Text(186.0, 1087.2, 'NO_2 <= 17.5\ngini = 0.204\nsamples = 17\nvalue = [23, 3]\nnclass
= Yes'),
Text(93.0, 776.5714285714287, 'gini = 0.0\nsamples = 11\nvalue = [17, 0]\nnclass = Ye
s'),
Text(279.0, 776.5714285714287, 'gini = 0.444\nsamples = 6\nvalue = [6, 3]\nnclass = Ye
s'),
Text(372.0, 1087.2, 'gini = 0.0\nsamples = 522\nvalue = [815, 0]\nnclass = Yes'),
Text(651.0, 1397.8285714285716, 'PM10 <= 4.5\ngini = 0.041\nsamples = 64\nvalue = [2,
94]\nnclass = No'),
Text(558.0, 1087.2, 'gini = 0.0\nsamples = 35\nvalue = [0, 56]\nnclass = No'),
Text(744.0, 1087.2, 'O_3 <= 52.0\ngini = 0.095\nsamples = 29\nvalue = [2, 38]\nnclass
= No'),
Text(651.0, 776.5714285714287, 'gini = 0.0\nsamples = 12\nvalue = [0, 19]\nnclass = N
o'),
Text(837.0, 776.5714285714287, 'NOx <= 23.5\ngini = 0.172\nsamples = 17\nvalue = [2,
19]\nnclass = No'),
Text(744.0, 465.9428571428573, 'gini = 0.298\nsamples = 10\nvalue = [2, 9]\nnclass = N
o'),
Text(930.0, 465.9428571428573, 'gini = 0.0\nsamples = 7\nvalue = [0, 10]\nnclass = N
o'),
Text(2406.375, 1708.457142857143, 'NOx <= 31.5\ngini = 0.402\nsamples = 1419\nvalue =
[628, 1628]\nnclass = No'),
Text(1581.0, 1397.8285714285716, 'NMHC <= 0.035\ngini = 0.409\nsamples = 94\nvalue =
[102, 41]\nnclass = Yes'),
Text(1488.0, 1087.2, 'NOx <= 25.5\ngini = 0.497\nsamples = 58\nvalue = [48, 41]\nclas
s = Yes'),
Text(1209.0, 776.5714285714287, 'CH4 <= 1.375\ngini = 0.43\nsamples = 22\nvalue = [1
0, 22]\nnclass = No'),
Text(1116.0, 465.9428571428573, 'gini = 0.0\nsamples = 6\nvalue = [9, 0]\nnclass = Ye
s'),
Text(1302.0, 465.9428571428573, 'SO_2 <= 2.0\ngini = 0.083\nsamples = 16\nvalue = [1,
22]\nnclass = No'),
Text(1209.0, 155.3142857142857, 'gini = 0.18\nsamples = 8\nvalue = [1, 9]\nnclass = N
o'),
Text(1395.0, 155.3142857142857, 'gini = 0.0\nsamples = 8\nvalue = [0, 13]\nnclass = N
o'),
Text(1767.0, 776.5714285714287, 'CH4 <= 1.395\ngini = 0.444\nsamples = 36\nvalue = [3
8, 19]\nnclass = Yes'),
Text(1674.0, 465.9428571428573, 'NO_2 <= 23.5\ngini = 0.05\nsamples = 24\nvalue = [3
8, 1]\nnclass = Yes'),
Text(1581.0, 155.3142857142857, 'gini = 0.0\nsamples = 19\nvalue = [34, 0]\nnclass = Y
es'),
Text(1767.0, 155.3142857142857, 'gini = 0.32\nsamples = 5\nvalue = [4, 1]\nnclass = Ye
s'),
Text(1860.0, 465.9428571428573, 'gini = 0.0\nsamples = 12\nvalue = [0, 18]\nnclass = N
o'),
Text(1674.0, 1087.2, 'gini = 0.0\nsamples = 36\nvalue = [54, 0]\nnclass = Yes'),
Text(3231.75, 1397.8285714285716, 'PM25 <= 10.5\ngini = 0.374\nsamples = 1325\nvalue
= [526, 1587]\nnclass = No'),
Text(2604.0, 1087.2, 'CO <= 0.25\ngini = 0.319\nsamples = 911\nvalue = [287, 1157]\nnc
lass = No'),
Text(2232.0, 776.5714285714287, 'PM25 <= 2.5\ngini = 0.426\nsamples = 87\nvalue = [9
2, 41]\nnclass = Yes'),
Text(2046.0, 465.9428571428573, 'CH4 <= 1.385\ngini = 0.5\nsamples = 32\nvalue = [25,
25]\nnclass = Yes'),

```

```
Text(1953.0, 155.3142857142857, 'gini = 0.083\nsamples = 14\nvalue = [22, 1]\nnclass = Yes'),
Text(2139.0, 155.3142857142857, 'gini = 0.198\nsamples = 18\nvalue = [3, 24]\nnclass = No'),
Text(2418.0, 465.9428571428573, 'NMHC <= 0.045\ngini = 0.311\nsamples = 55\nvalue = [67, 16]\nnclass = Yes'),
Text(2325.0, 155.3142857142857, 'gini = 0.435\nsamples = 36\nvalue = [34, 16]\nnclass = Yes'),
Text(2511.0, 155.3142857142857, 'gini = 0.0\nsamples = 19\nvalue = [33, 0]\nnclass = Yes'),
Text(2976.0, 776.5714285714287, 'station <= 28079016.0\ngini = 0.253\nsamples = 824\nvalue = [195, 1116]\nnclass = No'),
Text(2790.0, 465.9428571428573, 'CO <= 0.35\ngini = 0.137\nsamples = 762\nvalue = [89, 1114]\nnclass = No'),
Text(2697.0, 155.3142857142857, 'gini = 0.216\nsamples = 357\nvalue = [73, 520]\nnclass = No'),
Text(2883.0, 155.3142857142857, 'gini = 0.051\nsamples = 405\nvalue = [16, 594]\nnclass = No'),
Text(3162.0, 465.9428571428573, 'O_3 <= 19.5\ngini = 0.036\nsamples = 62\nvalue = [106, 2]\nnclass = Yes'),
Text(3069.0, 155.3142857142857, 'gini = 0.0\nsamples = 50\nvalue = [88, 0]\nnclass = Yes'),
Text(3255.0, 155.3142857142857, 'gini = 0.18\nsamples = 12\nvalue = [18, 2]\nnclass = Yes'),
Text(3859.5, 1087.2, 'EBE <= 0.25\ngini = 0.459\nsamples = 414\nvalue = [239, 430]\nnclass = No'),
Text(3627.0, 776.5714285714287, 'CH4 <= 1.395\ngini = 0.407\nsamples = 81\nvalue = [98, 39]\nnclass = Yes'),
Text(3534.0, 465.9428571428573, 'CH4 <= 1.355\ngini = 0.14\nsamples = 61\nvalue = [98, 8]\nnclass = Yes'),
Text(3441.0, 155.3142857142857, 'gini = 0.0\nsamples = 50\nvalue = [88, 0]\nnclass = Yes'),
Text(3627.0, 155.3142857142857, 'gini = 0.494\nsamples = 11\nvalue = [10, 8]\nnclass = Yes'),
Text(3720.0, 465.9428571428573, 'gini = 0.0\nsamples = 20\nvalue = [0, 31]\nnclass = No'),
Text(4092.0, 776.5714285714287, 'NOx <= 91.0\ngini = 0.39\nsamples = 333\nvalue = [141, 391]\nnclass = No'),
Text(3906.0, 465.9428571428573, 'station <= 28079016.0\ngini = 0.498\nsamples = 56\nvalue = [46, 40]\nnclass = Yes'),
Text(3813.0, 155.3142857142857, 'gini = 0.229\nsamples = 24\nvalue = [5, 33]\nnclass = No'),
Text(3999.0, 155.3142857142857, 'gini = 0.249\nsamples = 32\nvalue = [41, 7]\nnclass = Yes'),
Text(4278.0, 465.9428571428573, 'O_3 <= 1.5\ngini = 0.335\nsamples = 277\nvalue = [95, 351]\nnclass = No'),
Text(4185.0, 155.3142857142857, 'gini = 0.357\nsamples = 52\nvalue = [66, 20]\nnclass = Yes'),
Text(4371.0, 155.3142857142857, 'gini = 0.148\nsamples = 225\nvalue = [29, 331]\nnclass = No')]
```



```
In [74]: print("Linear:",lis)
print("Lasso:",las)
print("Ridge:",rrs)
print("ElasticNet:",ens)
print("Logistic:",los)
print("Random Forest:",rfcs)
```

Linear: 0.9996222720077558
 Lasso: 0.3728582371674869
 Ridge: 0.9979283454674264
 ElasticNet: 0.5889752318704515
 Logistic: 0.5536888239590942
 Random Forest: 0.9802696314989101

Best model is Linear Regression

In []: